

Parameter passing, and guaranteed unified initialization

Document Number: **d0708 R0 – draft 21**

Date: 2020-09-16

Reply-to: Herb Sutter (hsutter@microsoft.com)

Audience: EWG, LEWG

“How do I pass parameters?” and “how do I initialize variables?” empirically accounts for ~23% of the total rules in major C++ guideline literature, most of which would be eliminated by this proposal. This paper aims to:

Provide **declarative** parameter passing **consistently** for all parameters (templates and non-templates).

Reduce tedious boilerplate code, including to automate laborious rote **optimizations** currently done by hand.

Provide strong **initialization-before-use guarantees** uniformly for objects of all types (PODs and non-PODs).

Provide true **unified initialization** syntax and semantics.

Contents

0	Overview.....	2
0.1	Background and motivation.....	2
0.2	Root cause: “How” to pass parameters, by idiom and convention	2
0.3	Design principles.....	4
0.4	Acknowledgments.....	4
1	“What, not how” parameter passing	5
1.1	“What”: Directly declare intent	5
1.2	Definite first/last use (see also P1179)	8
1.3	Parameters.....	9
1.4	Range- for	15
1.5	Return values	16
1.6	Overloading.....	17
1.7	Segue: The this parameter	18
2	Guaranteed and unified initialization.....	19
2.1	out and consistent =uninitialized	19
2.2	out (including out this) and explicit named constructors	22
2.3	Consistent =initializer	23
2.4	Segue toward generalized operator= : Zero/multiple params.....	24
2.5	Generalized operator= , part 1: in params	25
2.6	Generalized operator= , part 2: out “this”	26
2.7	Together: Unified {copy,move}×{construction,assignment}	27
	Appendix: Implementation options for §1	29
	Basic option 1: Static.....	30
	Basic option 2: Dynamic.....	31
	Current selected approach: Static for forward , dynamic otherwise	32
	Bibliography	34

0 Overview

0.1 Background and motivation

C++ has perhaps the largest guidance and coding standard literature of any programming language. The bad news is that there is so much to teach; the good news is that our community has already documented what we think are the important things a C++ programmer has to know, in a way that can be measured and used to draw conclusions about how to improve the language.

I am performing a review of the C++ guidance literature to catalog and analyze all the rules they teach. Of the 638 rules I have catalogued and analyzed so far, ~16% are about “how do I pass parameters?” and another ~7% are about “how do I initialize variables?” (including ~2% of informational explanations such as “what if I don’t have program-meaningful values when I have to declare a variable?” and “is = in a declaration an assignment?”).

Most of those ~23% of all published C++ guidance would be eliminated by the features in this paper, so that we would not have to teach them anymore for new modern code. The resulting code will not only be simpler, but will be more robust and correct (by eliminating classes of errors), and at least as efficient (and often more efficient, by eliminating the need for useless work such as dead writes which are required today).

0.2 Root cause: “How” to pass parameters, by idiom and convention

In today’s C++, we require the programmer to specify the mechanics of “how” to pass a parameter, and so “what” the parameter is intended to be used for is expressed only by indirect convention. This created a need to invent (eventually multiple kinds of) references, particularly for natural operator overloading.

Today, a function can take a parameter via any combination of 10 major options:

$\{ \text{const, non-const} \} \times \{ \text{value, pointer, lvalue reference, rvalue reference, forwarding reference} \}$

However, some of these are **meaningless**, so we have to teach not to use them, and why. For example:

- An “rvalue reference to **const**” parameter is meaningless, because it prevents moving.
- A “forwarding reference to **const**” parameter is probably meaningless, because it would forward only lvalue/rvalue-ness while treating the object as const, and since “rvalue reference to **const**” is meaningless its only remaining useful meaning would be the same as just declaring an lvalue reference to **const**.

Worse, some combinations are **ambiguous**, so every major compiler and platform of the past 30 years has repeatedly reinvented nonstandard annotations to supply the missing information. For example:

- An **X&** parameter is ambiguous about whether the parameter is:
 - an “inout” parameter,
 - an initialized “out” parameter to be filled in by value, or
 - an uninitialized “out” parameter to be filled in by construction.

The usually correct guess is “inout” because that is the most common, but there is no way to express “out” which is important because the initial value should not be read.

- An **X*** parameter is ambiguous about whether the primary object being passed is:
 - the pointer itself (e.g., with C-style **FILE*** APIs the pointer itself is the value),

- an `X` that only for { flexibility, optionality } happens to be passed by pointer (which in turn may or may not be safe to store to outlive the function call),
- an initialized “out” parameter to be filled in by setting its value (usually assignment), or
- an uninitialized “out” parameter to be filled in by construction.

And so if the intent is one of the “out” cases we cannot enforce that the initial value should not be read.

- An `X**` parameter is ambiguous about whether it means:
 - an `X**` “in” parameter,
 - an `X*` “inout” parameter, or
 - an `X*` “out” parameter.

And, unfortunately, some combinations are useful but also **difficult or impossible to write**. For example:

- A “forwarding reference to `X`” parameter where `X` is a concrete type is difficult to write, though may be possible by writing the function instead as a template whose type parameter is constrained to `X` using `requires` or `enable_if`.
- An “rvalue reference to `T`” parameter where `T` is a template parameter type is probably impossible to write, because just writing `T&&` would be a forwarding reference instead and it’s unclear (and unteachable) whether the workaround of writing `std::forward` instead of `std::move` has the same effect.

The following table summarizes most of the rules we repeatedly teach today:

	What we teach today: “How” mechanics
In	Pass by value for “cheap to copy/move” types (incl. builtin types) Otherwise, pass by const X& + Overload non-templated rvalue reference X&& + std::move once to optimize rvalues except if <code>X</code> must be a type parameter, write templated forwarding reference X&& + enable_if/requires is_reference_v<X> and std::forward instead except consider passing <code>X</code> by value if it’s an “in+copy” parameter to a constructor
In-out	Pass by non-const X&
Out	Pass by non-const X& + nonstd annotations Can’t distinguish from in-out in the language Can’t enforce write-before-read or must-write
Move	Pass by non-templated rvalue reference X&& + std::move once except if <code>X</code> must be a type parameter, write templated forwarding reference X&& + enable_if/requires !is_reference_v<X> and std::forward instead
Forward	Pass by templated forwarding reference T&& + std::forward once and if we want only a concrete type <code>X</code> , add enable_if/requires is_convertible_v<T,X>

There’s a lot wrong here. Most embarrassing of all is that **“in,” the most common case, is the worst**: It’s bad enough that for “in” we teach writing an overload set instead of a single function (never mind that the suggested overload doesn’t even work for templates where `&&` means something else), and possibly overload a function template (even if the original function was not a template)... except possibly for constructors (not even mentioned here, but see [Sutter 2014]). But that’s not all: For a function with `N` “in” parameters that aren’t cheap to copy, following the guidance requires writing 2^N overloads for each combination of lvalue and rvalue arguments... and so some experts just write a perfect forwarder instead even though that is very advanced and isn’t even expressing “in” read-only-ness anymore... and we have failed.

This advice is essentially unteachable, which explains why it continues to be debated even among experts.

Note When move semantics were new, some trainers and style guides hoped that this would let us simplify the guidance to simply pass “in” objects by value always. That guidance definitely felt like a breath of fresh air compared to the complexity of teaching overloads and different kinds of references, and it did automatically move from rvalues, but it was flawed (led to usability and performance pitfalls) and it is not embraced by most guidance literature or the C++ standard library.

My CppCon 2014 talk [[Sutter 2014](#)] shows how complicated this gets even when seriously attempting to simplify the guidance. We’re still teaching and debating how to pass C++ parameters; there’s something wrong with that.

0.3 Design principles

Note These principles apply to all design efforts and aren’t specific to this paper. Please steal and reuse.

The primary design goal is conceptual integrity [[Brooks 1975](#)], which means that the design is coherent and reliably does what the user expects it to do. Conceptual integrity’s major supporting principles are:

- **Be consistent:** Don’t make similar things different, including in spelling, behavior, or capability. Don’t make different things appear similar when they have different behavior or capability. — For example, this proposal provides a single syntax for efficiently passing all “in” parameters, instead of teaching multiple divergent syntaxes for “in” parameters. This proposal also avoids today’s design problems of making `&&` reference parameters mean two incompatible things (rvalue reference and forwarding reference), making `&` reference parameters conflate two different meanings (`inout` and `out`), and having to explain that top-level `const` on a by-value parameter is meaningless except when it isn’t. It also allows using “`= initializer`” consistently for construction and assignment, and enables the guarantee that every variable is initialized before use with a program-meaningful value (never a dead write).
- **Be orthogonal:** Avoid arbitrary coupling. Let features be used freely in combination. — This proposal enables freely and conveniently composing `move` and `forward` with both template parameter types and concrete types equally, which allows declaring an “rvalue reference” parameter of a template parameter type and a “forwarding reference” parameter of a concrete type, both of which are difficult to express today. It also allows `=` to be used with the same meaning orthogonally to set an object’s value, whether the value is being set at declaration time or not, and whether it performs construction or assignment.
- **Be general:** Don’t restrict what is inherent. Don’t arbitrarily restrict a complete set of uses. Avoid special cases and partial features. — This proposal supports extending the parameter passing model generally throughout the language where an object is passed from one piece of code to another, including range-`for` loops. It also supports for the assignment operator all the options already allowed for constructors, including assignment from zero and multiple arguments, to remove C++’s current arbitrary restriction.

These also help satisfy the principles of least surprise and of including only what is essential, and result in features that are additive and so directly minimize concept count (and therefore also redundancy and clutter).

Additional design principles include: Make important things and differences visible. Make unimportant things and differences less visible. — This proposal makes `inout` and `out` distinct, and suggests considering making arguments taken by mutable reference more visible at call sites.

0.4 Acknowledgments

Thank you to the following for reviews and comments on this material: Joe Bialek, Marshall Clow, Gabriel Dos Reis, Timur Doumler, Lee Howes, Nicolai Josuttis, Thomas Köppe, Alisdair Meredith, Gor Nishanov, Barry Revzin, Nathan Sidwell, Michael Spertus, Bjarne Stroustrup, Andrew Sutton, Daveed Vandevoorde, Titus Winters.

1 “What, not how” parameter passing

1.1 “What”: Directly declare intent

“References were introduced primarily to support operator overloading. ... C passes every function argument by value, and where passing an object by value would be inefficient or inappropriate the user can pass a pointer. This strategy doesn’t work where operator overloading is used. In that case, notational convenience is essential...” — [Stroustrup94] p. 86

This paper proposes to upgrade our “how” convention to an abstraction: By enabling the programmer to directly declare their intent of “what” they want to use the parameter for, we can leave the optimal delivery mechanics to the language and machine target, improving performance and correctness as well as simplicity.

This paper proposes five explicit ways to pass parameters, roughly in order from most to least commonly used:

```
auto f(    in X x) { }      // present me an X I can read from
auto f(  inout X x) { }    // present me an X I can read from and will write to
auto f(    out X x) { }    // present me an X I will assign to

auto f(  move X x) { }     // present me an X I will move from
auto f(forward X x) { }    // present me an X I will pass along
```

Note Syntax is bikesheddable. This syntax is unambiguous when parameter names are required.

And it proposes two ways to return values:

```
auto f () move    X { }    // move an X to the caller
auto f ()         X { }    // (possibly same; we’re already partway to ‘move’ default)
auto f () forward X { }    // pass along an X to the caller
```

The need for declaring the data flow direction has been rediscovered many times, in different languages (e.g., [Ada](#) and [C#](#)) and in C and C++ commercial projects using ad-hoc convention or nonportable annotations for tools like DOxygen and SAL. As UX study participants were quick to point out, when they are annotations they always get out of sync, whereas when they are in the language they stay in sync and enable compiler optimizations.

1.1.1 Implementation: Delegation to platform ABI (see also Appendix)

The platform ABI specifies certain details of the argument passing convention. This paper describes the semantics, and the ABI specifies how those semantics are implemented. In some cases, the ABI has the option of specifying that the implementation be via an overload, or by passing additional bits in available CPU flags or combined in an additional register.

For example, for an `in X` parameter: If `X` is small and trivially copyable (e.g., `int`) it is passed by value, and the platform ABI defines what is the limit of “small” (e.g., 64 bits or smaller). Otherwise, `X` is passed by pointer with the semantics that rvalue arguments are automatically moved from, and the platform ABI defines whether that is accomplished by emitting overloads on `const&` and `&&` (as we teach people to today by hand, but which is combinatorial) or by passing an `is_rvalue` flag for the parameter that is filled by the caller (and if so, where the flag is stored, in a CPU flag or combined in an additional register with other parameters’ flags if there are any). If flags are passed in registers, the ABI documents how bits related to different parameters are combined in a bit-field to minimize stack frame overhead.

1.1.2 Summary of parameter passing semantics

The following table summarizes the semantics of each of the five declared intents. Overloading is allowed, and the compiler will select the best match based on the argument restrictions. The following sections go into more detail for each including examples.

Note that all of the following are compatible with calling existing functions. For example, an `in` parameter is simply treated as a `const` lvalue for all purposes including for passing it onward to an existing overload set, except that the definite last use treats it as an rvalue (only if the actual argument is an rvalue and the use would be legal for an rvalue) which naturally selects any rvalue overloads for the function it's being passed to.

For the last two, treating the parameter as a `const` lvalue for non-last uses allows the function to read the 'in' value (e.g., `forward` parameters commonly want to do trace/log output of the about-to-be forwarded value).

Highlights here and in §1.3.1 and §1.3.4 are speculative additions for possible extension to destructive move.

_____ X x	Calling convention	Caller arguments	Callee parameter uses
<code>in</code>	X (by value) if cheaply trivially copyable (defined by platform ABI) Otherwise, X*	Any initialized object (lvalue or rvalue)	x is treated as a <code>const</code> lvalue, except that if the argument is a non- <code>const</code> rvalue then each definite last use of x that would be legal for an rvalue treats x as a non- <code>const</code> rvalue (as-if casting it to an xvalue) (including can move from an rvalue argument)
<code>inout</code>	X*	Any initialized non- <code>const</code> lvalue	x is treated as a non- <code>const</code> lvalue If the function is nonvirtual, at least one path must contain a non- <code>const</code> use of x (otherwise, the parameter should be <code>in</code>)
<code>out</code>	X*	Any uninitialized or non- <code>const</code> lvalue	x is treated as a <code>const</code> lvalue, except that assignment-expression to x is also allowed Every path must contain a definite first use of x, and that definite first use must be either to pass x to another function's <code>out</code> parameter, or else be an assignment-expression to x (otherwise, the parameter should be <code>inout</code>) If the first use is an assignment-expression, then if the argument is uninitialized it constructs x with the right-hand side expression treated as the constructor argument list; otherwise, it performs an assignment as usual
<code>move</code>	X (by value) if cheaply trivially copyable (defined by platform ABI) Otherwise, X*	An initialized non- <code>const</code> rvalue or initialized definite last use lvalue	x is treated as a non- <code>const</code> lvalue, except that each definite last use of x treats x as a non- <code>const</code> rvalue (as-if casting it to an xvalue)
<code>forward</code>	X*	Any initialized object (lvalue or rvalue)	x is treated as a <code>const</code> lvalue, except that each definite last use of x preserves the argument's cv-qualification and value category

Notes No call site ever performs a non-trivial copy or move for an argument to these five parameter kinds.

Currently above, the “pass small trivially copyable types by value” optimization is only applied to `in` and `move`. We are considering generalizing it to perform this optimization uniformly in both directions which would additionally cover `inout`... the generalization for small trivially copyable types is: (a) pass all “in” by value (includes not only `in`, but also the “in” part of `inout`, and `move`) and return all “out” by value (includes `out` and the “out” part of `inout`). With this generalization, an `inout` parameter that is small and trivially copyable would not have to be passed by an indirection at all as it must today, but could be passed by copy in and by copy out which can be more efficient because then the `inout` parameter could be passed entirely in registers.

1.2 Definite first/last use (see also P1179)

This paper uses the following concepts, which is a strict subset of the [\[P1179R1\]](#) Lifetime rules. Other modern languages already successfully use similar path-based rules, such as C#'s definite assignment rules.

A “**local variable**” means any parameter, local variable, or temporary object in the scope of a function definition.

A function's “**control flow graph (CFG)**” has its usual meaning, except herein we ignore exceptional paths. A “**node**” has its usual meaning in a CFG except that each subexpression that mentions a local or nonlocal variable (including a local variable's declaration or destruction) is treated as a distinct node (rather than combined in larger basic blocks). A “**path**” is a connected directed sequenced-before path of nodes in a function's CFG.

“**Use**” of a local variable `var` means a node corresponding to a mention of `var`'s name that accesses `var`'s value (not its declaration, not just taking its address). A “**const use**” means a use that only reads `var`. A “**non-const use**” means a use that is not a `const` use (such as passing `var` to a function that takes it via reference to non-`const`). A “**definite first (or, last) use**” of `var` means a use that is not preceded (or, followed) by any other use of `var` on any path that shares that node, and is not in a loop body or on any path that traverses a backward `goto`.

1.2.1 Example

For example:

```
void f(int i, int j) {
    use(i);                // definite first use of i
    if (cond) {
        if (cond2)
            use(i,j);      // definite first use of j
    } else {
        use(i,j);          // definite first use of j
    }
    use(i);
    use(j);
    for (auto x : {1,2,3}) {
        if (cond3)
            use(x,i,j);
    }
    if (cond3)
        use(j);            // definite last use of j
}
```

Notes A local variable can have zero or more definite first uses, and zero or more definite last uses.

A variable may not have a definite first or last use on every path. In the above example, `i` does not have a definite last use on any path. Some kinds of parameter passing described in this document require that the parameter does have a definite first and/or last use on every path.

These rules will be simple for programmers to reason about, because they correspond directly to a simple linear visual scan of the code's nested blocks. C# programmers, who greatly value simplicity, have found C#'s “definite assignment” easy to reason about, and that feature's 12-page specification (see [\[ECMA-334\]](#), section 10.4.4, pp. 77-88) is significantly more complex than the above.

1.3 Parameters

1.3.1 in parameters

*“An internal memo dated January 1981 [Stroustrup,1981b] describes the idea: ‘Until now it has not been possible in C to specify that a data item should be read only, that is, that its value must remain unchanged. Neither has there been any way of restricting the use of arguments passed to a function. Dennis Ritchie pointed out that if **readonly** was a type operator, both facilities could be obtained easily...’ ” —B. Stroustrup (D&E, p. 89)*

```
auto f(in X x) { }           // present me an X I can read from
```

An **in** parameter is a parameter that the function body can read from. This case optimizes based on the size of the object and the lvalueness of the argument. For example:

C++20 equivalent	This paper (proposed)
<pre>void f1(const X& x) { g(x); } void f1(X&& x) { // overload to optimize rvalues g(std::move(x)); // remember to move only once }</pre>	<pre>void f1(in X x) { g(x); }</pre>
<pre>template<typename T> void f2(const T& t) { // can't easily overload for rvalues, & g(t); // nobody writes the requires/enable_if } // to pass builtins by value...</pre>	<pre>template<typename T> void f2(in T t) { g(t); }</pre>
<pre>template<typename T> bool should_pass_by_value_v = std::is_trivially_copyable_v<T> && sizeof(T) < _SOME_MAX; template<typename T> requires should_pass_by_value_v<T> void f3(T t) { my_trace("f", t); container.emplace_back(t); } template<typename T> requires (!should_pass_by_value_v<T>) void f3(const T& t) { my_trace("f", t); container.emplace_back(t); } template<typename T> requires (!should_pass_by_value_v<T> && !std::is_reference_v<T>) void f3(T&& t) { my_trace("f", t); // teaching: why no move here? container.emplace_back(std::forward<T>(t)); // erm, "move" }</pre>	<pre>template<typename T> void f3(in T t) { my_trace("f", t); container.emplace_back(t); }</pre>

Highlights here and in §1.1.2 and §1.3.4 are speculative additions for possible extension to destructive move.

For **in**, as also described for **out** (see §1.3.3), whether or not an exception is thrown, the function informs the caller whether an rvalue argument was moved from using a `[[relocates]]` move, and if so and the caller argument was a definite last use then the caller can elide its destructor.

Notes This lets us improve current C++ parameter passing guidance for “in” parameters in several ways.

First, it makes writing optimal code simple (and the default) for passing cheap-to-copy objects. In C++ we have always taught people to pass “small” or cheap-to-copy objects by value as `(X x)`, and “large” expensive-to-copy objects by reference as `(const X& x)`, and then teach what “small” means which changes over time on different platforms. Today’s guidance on this is not consistent; some authors teach “pass only built-ins by value,” others teach “pass built-ins or small-looking objects by value,” and both may be dated because actual measurements on modern platforms show that larger-than-historical objects are cheaper to pass by value which is another reason to leave it to the platform ABI. In this design, the programmer just passes `(in X x)` and the platform ABI specifies what “cheap to copy” means in a deterministic way for that platform ABI. Not only do programmers get to write simpler code and no longer have to remember this overloading-style guidance (at all and especially not per-platform), but they also get better portable performance as the same code is automatically implemented appropriately on each given platform.

Second, this is especially useful for templates. Essentially nobody tries to apply `requires` or `enable_if` and overloading to optimize cheap-to-copy `T`. – In this proposal, `in` does that naturally.

Third, it makes the “in+copy” case easier where the function will (or may) keep a copy of the parameter. In C++20, for this case we teach programmers who need to optimize for rvalues to overload on `(const X& x)` and copying from `x` in the body, and `(X&& x)` and writing `std::move(x)` in the body and on the last use only, with the function bodies usually otherwise identical. The C++20 standard library contains many examples of this duplication. – In this proposal, no hand-written overload is required, and in the body of the function we do not have to remember whether to use an explicit `std::move`, and call sites can correctly elide destructors after last use for moved-from objects when a relocating move is invoked (see §1.3.3).

Fourth, it avoids the long-standing wart in C and C++ that `f(T t)` and `f(const T t)` declare the same function. The confusion does not arise here because “in” parameters are never `const`-qualified and are always treated `const` in the function body.

Finally, it enables an additional optimization in environments that perform calls outside the program’s address space and so always copy data to/from the callee (e.g., RPC, sending data to GPU memory), because `in` lets the compiler automatically elide copying in one direction by construction.

There is one design point that we should validate with field experience: If an “in” parameter is aliased by another parameter or global, then although the `in` parameter is treated as `const` in that it cannot be modified directly, its value may change through the alias. This could be a pitfall, which we should determine by getting field usability experience. However, the [P1179R1] Lifetime rules default to banning passing non-owning Pointers that alias, so this may not be a problem in practice.

We can consider treating a local variable as an rvalue on definite last use, similar to an `in` parameter. Then examples like this would move:

```
void f(in X);

void g() {
    X x;
    f(x);           // definite last use of x, so treat as rvalue
}
```

```
void f(in X x) {
    X x2 = x;    // this is a move if the argument is an rvalue
}
```

1.3.2 inout parameters

```
auto f(inout X x) { }    // present me an X I can read from and will write to
```

An **inout** parameter is a parameter that the function body can read from and write to. For example:

C++20 equivalent	This paper (proposed)
<pre>void f1(<i>/*inout*/</i> X& x) { // cannot distinguish inout vs. out g(x); // ok x = 42; // ok but can omit }</pre>	<pre>void f1(inout X x) { use(x); // ok x = 42; // ok and required }</pre>
<pre>void f2(<i>/*inout*/</i> X& x) { // cannot distinguish inout vs. out y = x * 2; // ok } // not flagged: did not write to x</pre>	<pre>void f2(inout X x) { y = x * 2; // ok } // error, did not write to x</pre>

Note This is clearer than current C++ guidance. Today, both “in-out” and “out” are expressed by passing (**X& x**), so they are neither distinguished to the reader of the code nor enforceable by the language. Yet that is an important distinction to make. In this design, the programmer explicitly distinguishes “in-out” which allows reading from the parameter, and “out” which enforces that the parameter is not read until it is first written to.

1.3.3 out parameters

*“Finally, the [internal January 1981] memo introduces **writeonly**: ‘There is the type operator **writeonly**, which is used like **readonly**, but prevents reading rather than writing...’ “*
—B. Stroustrup (D&E, p. 90)

```
auto f(out X x) { }    // present me an X I will assign to
```

An **out** parameter is a parameter that the function body can read from only after writing to it, and will assign to it. (Note: The argument may be **uninitialized**; see §2.1.) For example, assuming **g(a)** has a non-**out** parameter:

C++20 equivalent	This paper (proposed)
<pre>void f1(<i>/*out*/</i> X& x) { // cannot distinguish inout vs. out g(x); // not flagged: read x = 42; // ok but can omit g(x); // ok }</pre>	<pre>void f1(out X x) { g(x); // error, trying to read “in” value x = 42; // ok, required g(x); // ok }</pre>
<pre>void f2(<i>/*out*/</i> X& x) { // cannot distinguish inout vs. out // ... no write to x ... } // not flagged: did not write to x</pre>	<pre>void f2(out X x) { // ... no write to x ... } // error, failed to write to x</pre>

If no exception is thrown, the function is guaranteed to have constructed an uninitialized argument. If an exception is thrown, if the argument was uninitialized it remains uninitialized (if it was initialized before the exception was thrown it is destroyed again in the callee).

For a parameter `out X x`, if the assignment-expression `x = expr;` is not a valid call to `operator=`, rewrite the call to `x = X{expr};` as a fallback.

1.3.4 move parameters

Highlights here and in §1.1.2 and §1.3.1 are speculative additions for possible extension to destructive move.

```
auto f(move X x) { }           // present me an X I will move from
```

A `move` parameter is a parameter that the function body will move from (consume the value of). For example:

C++20 equivalent	This paper (proposed)
<pre>void f1(X&& x) { container.emplace_back(std::move(x)); }</pre>	<pre>void f1(move X x) { container.emplace_back(x); }</pre>
<pre>template<typename T> // rref not allowed for templated T void f2(T&& t) { container.emplace_back(std::forward<T>(t)); } // + overload and =delete (const&) and (&) for lvalues</pre>	<pre>template<typename T> void f2 (move T t) { container.emplace_back(t); }</pre>

If no exception is thrown, the function is guaranteed to have moved from its argument. Whether or not an exception is thrown, the function informs the caller whether the argument was moved from using a `[[relocates]]` move, and if so and the caller argument was a definite last use then the caller can elide its destructor.

Notes This last paragraph is symmetric with `out`'s guarantees.

Many existing movable types, including `std::` smart pointers, already leave them in a non-owning state that does not require a destructor call (see P1029 for analysis). To be certain, however, we can require adorning a move operation with explicit `[[relocates]]` to designate a relocating move, which then triggers this optimization to automatically not destroy after a moved-from last use.

I'm currently naming this `move` for familiarity (and, unlike today, actually does guarantee a move), including in §1.3.6 regarding a potential call-site `move` qualifier instead of `std::move`. I think `consume` may be a better name, for the parameter and for any §1.3.6 call-site qualifier.

Declaring the parameter passing intent directly avoids the problems of expressing it indirectly today in terms of `const` and references, for example that we avoid today's problem that we have to explain why `const &&` rvalue references don't make sense, and `&&` rvalue references are onerous to write for templates.

This should eliminate the need to write `std::move` except when move will actually occur. As Bjarne Stroustrup notes, in today's teaching we recommend against the use of `std::move` just as we recommend against other casts; with this feature, we can do that more consistently. And when we do write `move`, this lets us teach "yes, `move` means `move`." (See also §1.3.6 regarding `move` at call sites.)

As Peter Dimov notes, we have a special rule for `return` (and, more recently, `throw`) that can be generalized so that the definite last use of an object is automatically made move-eligible. The problem is that we cannot easily do this in existing code without breakage; these new features enable us to do it without breakage because it applies to code that does not exist today. Please see [\[lib-ext 2017-02-08\]](#) and the followup replies.

Also, we no longer have to explain that “requests to move fall back to copy” with special rules. Instead, it’s just overload resolution on `in` and `move`: `in` can take any argument, and `move` takes only non-const rvalues, so if we have a call site that is `move` request and there is no `move` constructor in the overload set we naturally just select the `in` (copy) constructor by simple overload resolution, with no need for a special fallback rule.

1.3.5 forward parameters

```
auto f(forward X x) { }           // present me an X I will pass along
```

A `forward` parameter is a parameter that the function body will pass along to other code, and preserves the argument’s cv-qualification and value category. The function may forward a value of type `X` to other code (once), but not use the variable itself in any other way that could disturb it. The function’s body may read from `x` before it forwards it. For example:

C++20 equivalent	This paper (proposed)
<pre>template<typename T> void f1(T&& t) { container.emplace_back(std::forward<T>(t)); }</pre>	<pre>template<typename T> void f1(forward T t) { container.emplace_back(t); }</pre>
<pre>// difficult to write for non-templates template<typename widget> requires is_same_v<remove_cvref<T>, widget> void f2(widget&& w) { // move a concrete type ‘widget’ container.emplace_back(std::forward<widget>(w)); }</pre>	<pre>void f2(forward widget w) { container.emplace_back(w); }</pre>

Notes This allows a forwarding parameter without writing a template.

This also makes it easier to replace function-like macros. This macro:

```
#define F00(x,y) ((x) + (y))
```

can be replaced with the following function template in C++17:

```
template<typename X, typename Y>
auto Foo(X&& x, Y&& y) -> decltype(auto)
{ return std::forward<X>(x) + std::forward<Y>(y); }
```

and can be replaced with the following equivalent in C++20:

```
auto Foo(auto&& x, auto&& y) -> decltype(auto)
{ return std::forward<decltype(x)>(x) + std::forward<decltype(y)>(y); }
```

and can be replaced with the following equivalent in this proposal:

```
auto Foo(forward auto x, forward auto y) -> forward { return x + y; }
```

1.3.6 (optional) Call site explicit side effects

Optionally, we could also consider requiring that the call site be annotated in some way for every argument to an `inout`, `out`, `move`, or `forward` parameter, to avoid the common problem that the call site has no visual indication that the object it passes could be modified.

For example, given:

```
void f1(in      A a);
void f2(inout  B b);
void f3(out    C c);
void f4(move   D d);
void f5(forward E e);
```

We could default every argument to such a parameter to `const` unless explicitly annotated `out` (or possibly `mutable`) to opt into non-`const`, and we already have `std::move` (and could add `move`) to explicitly opt into rvalue-ness:

```
int main() {
    A mya; B myb; C myc; D myd; E mye;

    f1(mya);           // ok
    f2(myb);           // error: myb is treated as const by default
    f2(out myb);       // ok (or "f2(mutable myb)" or "f2(&myb)" etc.)
    f3(myc);           // error: myc is treated as const by default
    f3(out myc);       // ok (or "f3(mutable myc)" or "f3(&myc)" etc.)
```

For `move` parameters, we already require `move`:

```
f4(myd);              // error: myd is an lvalue
f4(std::move(myd));   // ok (already required in C++20)
f4(move myd);         // potential alternative
```

For `forward` parameters, which transparently forward onward to other code, this allows additional expressivity that's otherwise inconvenient to write:

```
f5(mye);              // ok: treated as const, forwards only value category
                        //      (f5 can only forward its argument to const functions)
f5(out mye);          // ok: forwards constness and value category
f5(move mye);         // ok: forwards as non-const rvalue
}
```

This trades off some inconvenience for clearer and more robust code. I see strong arguments on both sides.

Pro arguments:

- Addresses major classes of persistent real-world errors including exploited security vulnerabilities.

- Many C++ environments have reinvented these.
- Positive field experience in other widely used languages (e.g., C#).

Con arguments:

- “Verbosity.” “Nannying.”

1.4 Range-for

A range-`for` loop passes each object in the range as an argument to the loop body, and the loop iteration variable is already declared as a parameter. Applying the intentional parameter passing declares the loop’s intent:

```
for (    in X x : rng) { }           // an X I can read from
for (  inout X x : rng) { }         // an X I can read and write
for (    out X x : rng) { }         // an X I will assign to
for (  move X x : views::move(rng)) { } // an X I will move from (move iters required)
for (forward X x : rng) { }         // an X I will pass along
```

This enables declaratively self-documenting loops that directly declare things we cannot conveniently express in C++20, and so achieve better correctness and/or efficiency:

- Loops that read elements are correct and efficient, and convenient to write: A C++20 `for(auto x : rng)` loop copies each element which is a known performance pitfall for large objects, and allows modifications to `x` that will be silently thrown away because programmers rarely write `const` here to express a read-only loop. Alternatively, a C++20 `for(const auto& x : rng)` loop passes every element by reference even when copying would be more efficient. – In this proposal, `for(in auto x : rng)` automatically is both efficient and guaranteed to be read-only.
- Loops that modify elements can be stated explicitly and conveniently up front using `inout`.
- Loops that `move` from the elements in the range can be stated explicitly up front, and require `move_iterators`. In the foregoing, `views::move` is the range-v3 library’s name for a convenience wrapper that creates a pair of `move_iterators` from a range, the same way that `std::make_move_iterator` wraps an individual iterator.
- Loops that `forward` the elements in the range can be nongeneric: A C++20 `for(auto&& x : rng)` must be generic, whereas this design permits both the generic `for(forward x := rng)` and the nongeneric `for(forward x : X = rng)`.
- Loops can correctly initialize an array or other collection: A `for(out x := rng)` loop is guaranteed by construction to visit every element of `rng` and set it (construct or assign as appropriate) without reading the initial value (if any), and can deterministically initialize objects in an array.

1.5 Return values

These are primarily proposed for symmetry and to make current spellings more consistent, for example to spell the unique `->decltype(auto)` as a more regular `->forward`.

1.5.1 move return value

Given

```
auto f() -> move X { }      // move an X to the caller
```

we have a function that returns an rvalue of type `X`.

Note The default value return is already on a path to these semantics, including move from `return` of a local variable, and later move from `throw` of a local variable. If we like it, we could go further in the same directly by making this the whitespace default.

1.5.2 forward return value

A `forward` return value preserves the return-expression's cv-qualification and value category. A function definition is required.

Given

```
auto f() -> forward X { }    // pass along an X to the caller
```

we have a function that forwards a value of type `X`, preserving its cv-qualification and ref-qualification as deduced from the `return` statement. Explicit cv-qualification or ref-qualification is not allowed.

Note When returning an lvalue, this preserves the ability to naturally use lvalue returns in compound expressions (e.g., `v[0] = 42;`) including function chaining (e.g., `f(a).g(b).h(c);`).

A deduced return type defaults to `forward X` if all returns in the body return the same `forward` parameter of (possibly deduced) type `X`.

1.6 Overloading

Here is the matrix for viable arguments for each parameter type, according to whether the argument’s constness, l/r-valueness, and initialization state (initialization applies to local lvalues only).

	uninitialized	const lvalue	non-const lvalue	const rvalue	non-const rvalue
in	-	ok	ok	ok	ok
inout	-	-	ok	-	-
out	ok	-	ok	-	-
move	-	-	-	-	ok
forward	-	ok	ok	ok	ok

Note The “const rvalue” combination is included for completeness, ignoring the preexisting and independent question of whether that combination is actually useful (it arises mainly via `const X f();`).

Overloading on only these parameter passing options is not allowed. For example:

- Overloading `f(in X)` and `f(out Y)` is allowed, because they have different types.
- Overloading `f(int X)` and `f(out X, in Y)` is allowed, because they have different cardinality.
- Overloading `f(in X)` and `f(out X)` is not allowed, because they differ only in `in` vs. `out`.

Mixing these parameter passing options with today’s pass-by-value/`&/&&` is not allowed in an overload set of functions with the same name in the same scope.

Notes Overloading `in` and `forward` inherently makes call sites ambiguous and so cannot be allowed. Otherwise, all combinations have unambiguously better matches for any call site and so in principle “could” be allowed to overload. Similarly, mixing uses of these and the pass-by-value/`&/&&` styles “could” be allowed in an overload set because they are either naturally ambiguous or naturally unambiguous based on the constness and lvalue-ness of the call site’s arguments. But just because we can doesn’t mean that we should, because the goal is simplicity and to expose combinatorial expressiveness only when it is actually useful — so, “should” overloading these be allowed?

These parameter passing options now directly express the programmer’s intent of what the parameter is used for, including especially the direction of data flow and whether/what side effects are allowed on the argument. An overload set should have consistent semantics so that a call site can reason about the whole overload set’s meaning, including things like whether the argument can be modified (not `in`) or whether its ownership is transferred (`move`). So it makes much less sense to overload on only the parameter passing option, because if `f(t)` (including in a template) could have totally different effects depending on the value category of `t` then the call site can’t reason well about the meaning of the call.

Today, the most compelling and known-useful example of overloading these is to overload a function on `const&` (“in”) and `rvalue-&&` (“move”). However, the only reason we do that today is for the purpose of optimizing for rvalue arguments... which “in” now already does, which removes the known motivation for overloading “in” and “move.”

If we discover compelling use cases to overload these, we can easily relax this restriction in the future and allow the overloading, except only the only inherently ambiguous combination of `in` and `forward`.

1.7 Segue: The `this` parameter

For member functions, the `this` parameter is implicit and so can't be qualified using the same syntax as other parameters, and so the C++ convention is to express qualifiers on the `this` parameter via special additional grammar at the end of the member function parameter list. Declarative parameter passing for `this` would then go in the same place as all the other `this` qualifiers:

```
// qualifying "this" on member functions
auto f() in      { }          // present me an X I can read from
auto f() inout   { }          // present me an X I can read from and will write to
auto f() out     { }          // present me an X I will assign to
auto f() move    { }          // present me an X I will move from
auto f() forward { }          // present me an X I will pass along
```

Note how these correspond to, and subsume and extend, the existing member function qualifications:

Passing <code>this</code> object...	Similar to today's...	Notes
<code>auto X::f() in</code>	<code>auto X::f() const</code>	<code>in</code> additionally passes <code>*this</code> by value when that is cheaper
<code>auto X::f() inout</code>	<code>auto X::f()</code>	Today's existing "mutable" default
<code>auto X::f() out</code>	<code>X::X</code> (constructor)	See §2
<code>auto X::f() move</code>	<code>auto X::f() &&</code>	<code>move</code> additionally guarantees that if the function returns normally, <code>this</code> object is known to be moved-from
<code>auto X::f() forward</code>	n/a	Today there is no way to forward <code>this</code> object

2 Guaranteed and unified initialization

2.1 out and consistent =uninitialized

As a related feature, a variable of any type may be explicitly uninitialized by declaring it `= uninitialized` with the simple semantics of deferring the constructor call to the definite first use, which directly helps examples where a sensible program-meaningful initial value is not known at the point the variable must be declared:

- A declaration of a local variable `x` that is initialized with `= uninitialized`; allocates stack memory as usual but does not invoke a constructor.
- Every path containing `x`'s declaration must contain a definite first use of `x`, and that definite first use must be either to pass `x` to another function's `out` parameter, or else be an assignment-expression to `x` (which can be from a single-object expression or a possibly-empty `initializer_list`) that performs a single constructor call (usually direct-initialization) of `x` with the right-hand side expression treated as the argument to `x`'s constructor.

Note For security applications, the initialization can also set the padding bits.

- If two local variables `x` and `y` are explicitly uninitialized, then on every path `x`'s definite first use must precede `y`'s definite first use.

For example:

```
int i;                // i is uninitialized (current language rule)
int j = uninitialized; // same, except guarantees j will be initialized
if (cond) {           // allows true alternate initialization
    i = 1;             // i is initialized
    j = 2;             // j is initialized
} else {
    i = 100;           // i is initialized
    j = 200;           // j is initialized
}

std::string s = uninitialized; // s is uninitialized, no constructor is called
if (cond2) {
    s = "xyzyzy";        // s is constructed, calls std::string(const char*)
} else {
    // calls: fill_string(int val, out string s)
    fill_string(i+j, s);  // s is constructed, function will construct string
}
```

This guarantees initialization without having to eliminate dead writes for overwritten initial values. For example:

```
void fill_struct(out BigStruct s) { // guaranteed initialization
    // computation, then the definite first use of s is:
    s = computed_value;            // calls s's constructor
    // any const or non-const use of s is fine after this point
}
```

```
int f() {
    BigStruct data = uninitialized;           // no dead writes to eliminate
    do {
        fill_struct(data);                   // guaranteed to construct on first call, then
                                              // guaranteed to assign on each subsequent call
    } while (something);
}
```

Notes The output parameters of functions like `std::memset` and `std::copy` and similar functions should be treated as if declared with `out` parameters.

We still teach “declare variables as locally as possible” and “prefer to initialize variables at their point of declaration” most of the time; this proposal just also acknowledges the reality of situations where that advice is not appropriate, such as the examples later in this Note. Importantly, initialization of variables at their point of declaration continues to be implicitly encouraged with a major carrot: The language gives automatic type deduction support for the initialize-at-declaration case, whereas using `=uninitialized` requires naming a concrete (non-deduced) type.

Allowing an assignment-expression to invoke a constructor, and allowing an `out` parameter to call a constructor, is intentionally a step toward actual uniform initialization in C++ where all value-setting (construction and assignment) is spelled using “`= val`” syntax in all cases.

This lets us provide initialize-before-use guarantees, and also improve parameter passing guidance. First, it improves the correctness and efficiency of initializing large PODs, including arrays of PODs. Even though many coding styles encourage initialization-at-declaration, they universally make an exception for cases like (C++20 style)

```
db_format data;           // POD matching a data layout structure
read_next_chunk(&data);    // e.g., from disk or network
while (...) read_next_chunk(&data);
```

where requiring a redundant artificial initial value on the declaration is **unuseful** (because the value will be immediately overwritten), **expensive** (because compilers cannot always eliminate the dead write), and **error-prone** (because the language provides no way to prevent `read_next_chunk` from reading from uninitialized data, and to require that it does write to at least some part of the data). – In this proposal, we allow functions like `read_next_chunk` to express that they guarantee to initialize (write to) the parameter, guarantee that they will not read from it, and enable strong initialization guarantees without requiring artificial initial values.

Changing the `read_next_chunk` to return by value instead is also not cost-free as it requires at least a move, and is generally a bad design when it is called repeatedly because the alternative to using a caller-provided buffer is to perform an allocation on every call or to resort to brittle workarounds like a static buffer.

Note that the “out” pattern is not limited to POD buffers, but comes up regularly in well-written code like (C++20 style)

```
vector<byte> data;         // nicely encapsulated non-POD type
read_next_chunk(data);
while (...) read_next_chunk(data);
```

where it is key that the function reuse a caller’s storage instead of managing its own; this is what takes this kind of loop from (possibly large) N allocations down to a single allocation for the entire loop. The implementation of `read_next_chunk` does not need to perform any `const` operation on `data`; it will either assign to `data`, or `.clear` then `.push_back`, or `.resize(n)` then write through `[]` subscripting, or otherwise fill `data` using any other combination of non-`const` operations. Under the proposed design, this code and all of those techniques still work, but we can gain the guarantee that `read_next_chunk` does not read stale `data` and will actually store new `data` (important to avoid reading garbage data, or sensitive data).

Further, these are semantics you want for a potentially moved-from argument: If the caller passes a moved-from object, `out` expresses nearly all of the safe uses of that object, and no unsafe ones.

To distinguish initialized vs. uninitialized arguments, overload on `out` and `inout`.

For trivially copyable types, there is an additional optimization in environments that perform calls outside the program’s address space and so always copy data to/from the (e.g., RPC, sending data to GPU memory), because `out` lets the compiler automatically elide the first copy (to the callee’s memory) because by construction no data is to be copied in.

Use `inout` to express “maybe-`out`” parameters that are not read but are written to, such an `error_code` parameter that is only set on error paths. The only additional thing that a `maybe_out` parameter style could provide would be to statically enforce that the callee doesn’t read the input value (we could require that a definite first use must be at least a non-`const` operation, or even more strictly an assignment). At this point that one difference from `inout` doesn’t seem to warrant adding a separate option, so I’m considering it as just one of the uses of `inout`; but we can always add it as a separate option anytime if we decide that it’s important.

If the function throws an exception, it informs the caller whether an uninitialized `out` parameter was initialized.

If `X` is an array type, initializing every element of the array is required. However, the cost of doing so is greatly mitigated compared to today’s C++ because we never initialize dummy values that will only be overwritten with “real” program-meaningful values.

In summary, `=uninitialized` together with `out` parameters gives us consistent initialization guarantees for PODs and non-PODs. Declaring `=uninitialized` has the same meaning as the default initialization for stack-based POD variables, but additionally:

- works consistently for **all types** (not just PODs);
- **statically guarantees initialization** on definite first use, including via an initializing function call via `out` parameter (a common, even pervasive, style of initialization with program-meaningful values);
- **enables initialization to always use program-meaningful values**, never requiring a placeholder (including default-constructed) value that is just injecting dead writes intended to be overwritten;
- guarantees that the later initialization is **consistently spelled as `=initializer`**;
- guarantees that the later initialization is an efficient **single constructor call**; and
- last but not least, lets us give **consistent guidance** for all types (instead of routinely giving different guidance for large PODs and arrays, which today we routinely teach not to initialize on declaration) and starts to put us on a path to a consistent spelling of initialization.

2.2 out (including out this) and explicit named constructors

Because any function can initialize its `out` argument, any function with an `out` is effectively a named constructor. When the function is a non-member non-friend, it is always a ‘delegating’ constructor, that is, one that’s written in terms of other constructors. When `this` is `out`, as in the following example, the named constructor has the usual private access.

Being able to conveniently give meaningful names to constructors naturally makes them both

- “explicit,” because they are always a named function call; and
- “tagged,” because they can be given distinct names;

and so subsumes those two usual workarounds for implicit or ambiguous conversions with a direct way to express the intent that also makes the resulting calling code clearer and self-documenting.

For example, we could express the `vector` constructors from (count,value), (values...), and (first,last) unambiguously, similar to the `assign` overloads but able to act as constructors and using non-overloaded names to demonstrate how that would look:

```
template<class T, class Allocator>
class vector {
    // ...

public:
    void set_fill ( in size_type count, in T value,      in Allocator alloc = Allocator() ) out;
    void set_values( in initializer_list<T> values,      in Allocator alloc = Allocator() ) out;

    template<class InputIt>
    void set_from ( in InputIt first, in InputIt last, in Allocator alloc = Allocator() ) out;
    // ...
};
```

Even for `vector<int>`, the canonical confusing example, this makes calling code’s constructors unambiguous and clear, without having to resort to using `{ }` sometimes and teaching the differences between `()` and `{ }` including why they can give nonintuitive results):

```
vector v<int> = uninitialized;           // v is uninitialized, no constructor is called
if (cond) {
    v.set_fill(3, 42);                  // v is constructed with values { 42, 42, 42 }
} else if (cond2) {
    v.set_values(3, 42);                 // v is constructed with values { 3, 42 }
} else {
    v.set_from(myset.begin(), myset.end()); // v is constructed
}
```

Note that, as in the previous section, `out` continues to give us the ability to directly call alternative constructors for a locally declared object.

And because `out` parameters can be written on non-member non-friend functions, users can write their own constructors that delegate to other constructors.

2.3 Consistent `=initializer`

Building on the above consistent uninitialization via `=uninitialized`, next we add consistent initialization via `=initializer`, and furthermore consistent value-setting for both initialization and assignment:

- works consistently for **all types** (not just PODs);
- works consistently in **declaration and assignment** syntaxes;
- allows all value-setting to be **consistently spelled as `=initializer`**;
- last but not least, lets us give **consistent guidance** for all value-setting in C++.

We allow

```
MyType var = initializer;
```

and

```
MyType var = uninitialized;
// ...
var = initializer;           // definite first use ⇒ construction
// ...
var = initializer;           // otherwise, assignment
```

where `initializer` may be:

- a single identifier or expression;
- an *initializer-list*; or
- a parenthesized list that is a parameter list to a constructor (possibly `operator=`, see §2.4ff).

and when the destination object is already initialized, the semantics for assignment are that an `operator=` taking *initializer* is selected if available, otherwise a value is constructed from *initializer* and move-assigned.

In each of the following examples, the first two sets of lines have identical semantics (i.e., call the identical functions with identical parameters). The third line demonstrates that the identical syntax works for assignment.

Value-setting from a single variable or expression is pretty much what we already have, expanded as above to allow deferring a constructor call via `=uninitialized`:

```
string s1 = other;           // construct from other
string s2 = uninitialized;
s2 = other;                  // construct from other
s2 = other;                  // assign from other
```

Value-setting from an initializer list is pretty much what we already have, expanded again with `=uninitialized`:

```
vector<int> v1 = {1, 2, 3};   // construct from {1, 2, 3}
vector<int> v2 = uninitialized;
v2 = {1, 2, 3};              // construct from {1, 2, 3}
v2 = {1, 2, 3};              // assign from {1, 2, 3}
```

I propose also allow value-setting from a parenthesized list of arguments:

```
vector<int> v = (1, 2);           // construct from (1, 2)
vector<int> v = uninitialized;
v = (1, 2);                     // construct from (1, 2)
v = (1, 2);                     // assign from (1, 2)
```

Note that this proposes a minor breaking change, only after the token sequence `= (`, because today `(1, 2)` is valid using the comma operator. Such code is rare, and in common cases where the expressions have no side effects, such as just naming two variables or literals, today it is likely a latent bug that does not mean what the programmer intended. I argue that making `=` consistent outweighs backward compatibility with rare edge cases, and outweighs the consistency of having `(a,b,c)` mean the comma operator elsewhere but a list of values if preceded by `=`. Also, code that wants to maintain the current meaning can add another pair of `()`.

Note also that this enables us to express something we could never (easily) express before, namely “**default assignment**” to set an already-constructed value to a default state, which is asymmetrically missing in current C++ (allowed for construction but not for assignment):

```
vector<int> v = ();              // default construction
vector<int> v = uninitialized;
v = ();                         // default construction
v = ();                         // default assignment (new: currently not allowed)
```

For backward compatibility with C++20 types that do not have the nullary assignment operator described in the next section, we can make this mean to invoke the default constructor followed by the move constructor.

2.4 Segue toward generalized `operator=`: Zero/multiple params

At this point, the reader might anticipate that this paper would additionally propose that we allow an assignment operator to be declared with no parameters or multiple parameters, as we already do for constructors, for example:

```
class X {
    // ...
public:
    X& operator=();              // (nullary) default assignment operator
    X& operator=(int, string);   // assignment operator from (int, string)
};
```

and then by the prior rule in this section to prefer an assignment operator if available that will be preferred if present. This makes `operator=` consistent with constructors.

Yes, this paper does propose that. But we can be even more general (and efficient): Consider what happens if those are `in` parameters...

2.5 Generalized operator=, part 1: in params

`in` (§1.3.1) naturally gives the efficiency of copying from cheap-to-copy values without writing a by-value parameter, and otherwise from automatically moving from rvalue arguments without writing `const&/&&` overloads:

```
class X {
    // ...
public:
    X& operator=(in int i, in string s);    // i: by-value copy semantics
};                                         // s: by-reference as if const&/&& overloaded
```

What if the parameter of a copy constructor or copy assignment operator is declared `in`, which means “by value if cheap to copy, else as if with a `const&/&&` overload”? First, consider the cheap-to-copy case:

```
class Point {
    int x, y;
public:
    Point(in Point that);                // that: by-value copy semantics
    Point& operator=(in Point that);      // that: by-value copy semantics
};
```

This performs efficient copying using the normal ‘pass cheap-to-copy objects by value’ rule we teach in general, but that until now we have not allowed for copying special member functions.

Second, and more importantly, note these also naturally move from rvalues when the type is not cheap to copy:

```
class X {
    vector<widget> vw;
public:
    X(in X that);                        // that: by-reference as if const&/&& overloaded
    X& operator=(in X that);             // that: by-reference as if const&/&& overloaded
};
```

This makes them naturally **unified copy and move** operations. This can make code clearer by helping to prevent needless code duplication between copy and move functions that otherwise have the same structure (read from source member values the same way), which not only keeps copy and move consistent but makes them easier to write and maintain. Here they are side by side:

C++20 equivalent	This paper (proposed)
<pre>class X { // ... public: X(const X& that); X(X&& that); X& operator=(const X& that); X& operator=(X&& that); };</pre>	<pre>class X { // ... public: X(in X that); X& operator=(in X that); };</pre>

Of course, if copy and move naturally perform different work, they can be written separately as usual. Note that the following are the new ways to declare separate copy and move functions:

```

class X {
    vector<widget> vw;
public:
    X(in X that);           // copy (could be copy and move if it were alone)
    X(move X that);         // move (explicit overload for move if desired)

    X& operator=(in X that); // copy (could be copy and move if it were alone)
    X& operator=(move X that); // move (explicit overload for move if desired)
};

```

Note that for `operator=`, like all member functions, the current default for `this` is non-const \Rightarrow inout.

Next, consider what happens when we make `this` an out parameter.

2.6 Generalized `operator=`, part 2: out “this”

We already saw that this proposal allows `this` to be an out parameter, with the semantics that when invoked with an uninitialized object the members are treated as if `=uninitialized` and therefore naturally must be initialized, in order, in the body (i.e., an out “this” function is inherently (also) a constructor):

```

class X {
    widget w;
    gadget g;
public:
    X& operator=(*...) out {
        w = /*...*/;
        g = /*...*/;           // note: must init g after w (compiler enforced)
        return *this;
    }
};

```

This is a **unified construction and assignment** function. For example:

```

X x1 = /*...*/;           // calls operator=, constructs x.w and x.g
X x2 = uninitialized;
x2 = /*...*/;             // calls operator=, constructs x.w and x.g
x2 = /*...*/;             // calls operator=, assigns x.w and x.g

```

Note that the same `operator=` function is being called in each case. Of course, if construction and assignment naturally perform different work, they can be written separately as usual.

Note This function can behave as a constructor, and has the same efficiency as if using a mem-init-list, but no mem-init-list is required. Furthermore, this allows the flexibility of doing arbitrary work as needed before constructing a data member, including initializing data members along different paths as shown in §2.1. if that is useful, instead of being forced to first construct a data member with a dead-write placeholder value before we can calculate a real program-meaningful value in the body. So this allows strictly more flexibility and expressiveness (and often efficiency) than constructors, with no special-purpose syntax unique to constructors.

This can make code clearer by helping to prevent needless code duplication between constructors and assignment operators that initialize their members the same way, which not only keeps construction and assignment consistent but makes them easier to write and maintain. Of course, if construction and assignment naturally perform different work, they can be written separately as usual.

Side by side:

C++20 equivalent	This paper (proposed)
<pre>class X { // ... public: X(/*...*/); X& operator=(/*...*/); };</pre>	<pre>class X { // ... public: X& operator=(/*...*/) out; };</pre>

2.7 Together: Unified {copy,move}×{construction,assignment}

Finally, if we write `operator=` with both `out` “this” and taking a single `in` parameter of its own type we get a **single function that performs copy construction, move construction, copy assignment, and move assignment**:

```
class X {
    widget w;
    gadget g;
public:
    X& operator=(in X that) out { // constructs *this if uninitialized, else assigns
        w = that.w;              // moves from that if rvalue, else copies
        g = that.g;
        return *this;
    }
};
```

Side by side:

C++20 equivalent	This paper (proposed)
<pre>class X { // ... public: X(const X& that); X(X&& that); X& operator=(const X& that); X& operator=(X&& that); };</pre>	<pre>class X { // ... public: X& operator=(in X that) out; };</pre>

Of course, we can still write our move-specific and/or assignment-specific overloads if we want, if they can do work more efficiently using a different body:

```
// Consistent to teach: “out this” ⇒ constructor, “move that” ⇒ move operation
class X {
    widget w;
    gadget g;
public:
    void operator=(in X that) out;           // copy constructor
    void operator=(move X that) out;         // move constructor
    X& operator=(in X that);                 // copy assignment
    X& operator=(move X that);               // move assignment
};
```

And this is (in my opinion, wonderfully) consistent, because now we can really reach that “= means setting the value” consistently throughout the language:

- **A class author can write `operator=` for all value setting**, both construction and assignment. And it’s simple: By default they can conveniently write a single `=` function to enable copy construction, copy assignment, move construction, and move assignment, including naturally efficient move-from-rvalue semantics, but always with the option of writing separate functions if desired. In this proposal, there is now never a need to write a traditional constructor with its magic rules (e.g., can’t have a return value, requires unique mem-init-list syntax), because every constructor can be written `operator=(*...*) out` with (a) equal or better efficiency in all cases, and (b) strictly better initialization flexibility and initialization guarantees.
- **A class user can write `=` for all value setting**, both construction and assignment. We can teach users simply that “= means =” — all operations invoked using `=` syntax (both construction and assignment) can invoke an `=` function in the class. There is no need to teach that “= in a declaration doesn’t mean assignment and doesn’t call `operator=`” because now it can.

Note It is high time for consistent initialization and assignment in C++. With this extension, we can (finally) teach a simple rule for value-setting in C++: Always write `= initializer`, it is as efficient on construction as if you had omitted the `=`, and it is consistent with assignment.

Consistent initialization and assignment has long been a thorny problem in C++, and many experts have attempted to improve the situation. The three features of (1) parameter passing, (2) initialization, and (3) generalized value setting are all interrelated, and in this proposal embracing

(a) the path-based concept of definite first use, and

(b) `out` parameters to let that first use extend through function calls

constitutes a fundamental (and, I argue, necessary and timely) change in approach that directly enables expressing this clean unification also of value-setting.

If we pursue this direction, we would obsolete ~23% of all published C++ “gotcha” guidance and may put at risk our language’s hard-won reputation for arcane difficulty. But C++ still has other corner cases, so all is not lost; this paper only removes two of the largest single sources of C++ usability complexity and teaching effort in practice. Removing another 65% is the topic of other/future paper(s).

Appendix: Implementation options for §1

Most of the proposed declarative parameter passing conventions’ semantics require additional information to be communicated from the caller to the callee:

<code>___ X x</code>	In: Extra argument information (conceptual) Always known statically, at compile time
<code>in</code>	<code>bool __x_arg_is_const_or_lvalue</code>
<code>inout</code>	—
<code>out</code>	<code>bool __x_arg_is_uninitialized</code>
<code>move</code>	—
<code>forward</code>	<code>bool __x_arg_is_const</code> <code>bool __x_arg_is_lvalue</code>

Note The directions of the boolean flag meanings for `in` and `out` are chosen so that they are `false` in the common case, so that FFIs from other languages can just correctly pass `false/0`-bits by default.

Because the “in” information is known at compile time, we can communicate it statically or dynamically. This section describes those two major implementation options using the following base example. To show the two forms of `in`, `A` is small and trivially copyable and `B` isn’t.

```
void f(in A a, in B b, out C c, forward D d) {
    // code section 1
    use(a);
    // code section 2
    use(b);
    // code section 3
    c = {1, 2, 3};
    // code section 4
    use(d);
    // code section 5
}
```

Notes:

- These are all concrete types for clarity. The implementation strategies are identical in a template.
- For `out` called with an uninitialized argument, if the callee constructed the argument then in the case where the callee exits via an exception the argument must be destroyed. It is cleanest to do this in the callee, so that’s what the following code samples will show. An alternative, but *not* recommended, is to do the destruction in the caller, which would require extending the calling convention to return an additional `bool` whether the argument was constructed that each call site must test and conditionally perform the destruction, which is chatty and more complex and proliferates the check at each call site.

Basic option 1: Static

To pass the information statically, the function would be compiled as if it had been written as:

```
template< bool __b_arg_is_nonconst_rvalue,
          bool __c_arg_is_uninitialized,
          bool __d_arg_is_const,
          bool __d_arg_is_lvalue >
void __generated_f( A a, const B* b, C* c, D* d )
try {
    bool __c_arg_was_constructed_by_f = false;
    // code section 1
    use(a);                                     // unchanged
    // code section 2
    if constexpr (__b_arg_is_nonconst_rvalue)
        use(std::move(*b));                   // treat as non-const rvalue
    else
        use(*b);                               // treat as const lvalue
    // code section 3
    if constexpr (__c_arg_is_uninitialized)
        new (c) C{1, 2, 3} , __c_arg_was_constructed_by_f = true;
    else
        *c = {1, 2, 3};                       // or, if unavailable, c = C{1, 2, 3}
    // code section 4
    if constexpr (__d_arg_is_const)
        if constexpr (__d_arg_is_lvalue)
            use(const_cast<const D*>(*d));     // treat as const lvalue
        else
            use(std::move(const_cast<const D*>(*d))); // treat as const rvalue (weird?)
    else
        if constexpr (__d_arg_is_lvalue)
            use(*d);                           // treat as non-const lvalue
        else
            use(std::move(*d));                 // treat as non-const rvalue
    // code section 5
} catch(...) {
    if (__c_arg_was_constructed_by_f) c->~C();
}
```

Notes Using `if constexpr` is more convenient and less combinatorial than spelling out specializations.

These could be `constexpr` parameters as proposed in [\[P1045R1\]](#), and be operationally equivalent.

Basic option 2: Dynamic

To pass the information dynamically, the function would be compiled as if it had been written as:

```
void __generated_f( A a, const B* b, C* c, D* d,
                  bool __b_arg_is_nonconst_rvalue,
                  bool __c_arg_is_uninitialized,
                  bool __d_arg_is_const,
                  bool __d_arg_is_lvalue )
try {
    bool __c_arg_was_constructed_by_f = false;
    // code section 1
    use(a); // unchanged
    // code section 2
    if (__b_arg_is_nonconst_rvalue)
        use(std::move(*b)); // treat as non-const rvalue
    else
        use(*b); // treat as const lvalue
    // code section 3
    if (__c_arg_is_uninitialized)
        new (c) C{1, 2, 3} , __c_arg_was_constructed_by_f = true;
    else
        *c = {1, 2, 3}; // or, if unavailable, c = C{1, 2, 3}
    // code section 4
    if (__d_arg_is_const)
        if (__d_arg_is_lvalue)
            use(const_cast<const D*>(*d)); // treat as const lvalue
        else
            use(std::move(const_cast<const D*>(*d))); // treat as const rvalue (weird?)
    else
        if (__d_arg_is_lvalue)
            use(*d); // treat as non-const lvalue
        else
            use(std::move(*d)); // treat as non-const rvalue
    // code section 5
} catch(...) {
    if (__c_arg_was_constructed_by_f) c->~C();
}
```

Note These can be packed into bitflags to make the stack frame smaller, which would be operationally equivalent. We should just measure which is faster, a slightly smaller stack frame or fewer bit-masking instructions. In this example, we could pass a single additional parameter of type:

```
enum __f_param_flags {
    __b_arg_is_nonconst_rvalue = 1,
```

```

__c_arg_is_uninitialized = 2,
__d_arg_is_const = 4,
__d_arg_is_lvalue = 8
};

```

Current selected approach: Static for `forward`, dynamic otherwise

This table summarizes the pros and cons.

	Option 1: Static	Option 2: Dynamic
Pros	No run-time overhead, including no branches	<p>Can be an ordinary function: If <code>&f</code> works without this parameter passing style, it still works</p> <p>Minimum possible binary code space overhead: in many cases, even better than we can do right now by hand because today we must write overloads</p>
Cons	<p>Binary code space overhead, including code explosion: All the common <code>// code sections</code> are duplicated a combinatorial number of times (modulo COMDAT folding of identical bodies)</p> <p>Effectively every function with an <code>in</code>, <code>out</code>, or <code>forward</code> param becomes a template; and because <code>in</code> is the dominant case, this means most functions that use the new style will effectively become templates</p> <p>Can't take <code>&f</code></p>	<p>1 hinted branch per definite last use of an <code>in</code> param</p> <p>1 branch per definite first use of an <code>out</code> param</p> <p>3 branches per <code>forward</code> param</p> <p>However, for using multiple of these parameters in the same definite last use function call expression causes a combinatorial explosion</p> <p>Especially, multiple <code>forward</code> parameters are expected to be common, because of forwarded parameter packs</p> <p>Changes the calling convention of the function, <code>&f</code> gives a different function types</p>

Our current decision for prototyping is to use a variant of Option 1 (static) for `forward` (basically just directly automate today's forwarding pattern), and a variant of Option 2 (dynamic) for everything else.

This has a user-visible semantic design consequence: We teach that we can't take the address of an otherwise-ordinary function that has `forward` parameters; it is an implicit template. We also encourage forwarders to be small functions since they will be stamped out for each combination of parameter value category and constness. We think this is likely a sweet spot because forwarders already behave this way, and `forward` still gives benefit because it is strictly more convenient.

So the function would be compiled as if it had been written as:


```

template <class D>
void __generated_f( A a, const B* b, C* c, D&& d,
                  enum { __b_arg_is_nonconst_rvalue = 1,
                        __c_arg_is_uninitialized = 2 } __flags = 0 )
try {
    bool __c_arg_was_constructed_by_f = false;
    // code section 1
    use(a);                                // unchanged
    // code section 2
    if (__flags & __b_arg_is_nonconst_rvalue) [[unlikely]]
        use(std::move(*b));                // treat as non-const rvalue
    else
        use(*b);                           // treat as const lvalue
    // code section 3
    if (flags & __c_arg_is_uninitialized) [[unlikely]]
        new (c) C{1, 2, 3} , __c_arg_was_constructed_by_f = true;
    else
        *c = {1, 2, 3};                    // or, if unavailable, c = C{1, 2, 3}
    // code section 4
    use( std::forward<D>(d) );              // forwarding pattern
    // code section 5
} catch(...) {
    if (__c_arg_was_constructed_by_f) c->~C();
}

```

Because the `__flags` parameter defaults to 0, this is equivalent to providing an additional overload without any flag parameters at all that is friendly to non-C++ FFI callers. For example, developers who write this:

```
void f(in A a, in B b, inout C c, out D d)
```

would get this:

```

void f( A a, const B* b, [[inout]] C* c, [[out]] D* d,
      enum { __b_arg_is_nonconst_rvalue = 1,
            __d_arg_is_uninitialized = 2 } __flags = 0 )

```

and also this too which is foreign language-friendly:

```
void f( A a, const B* b, [[inout]] C* c, [[out]] D* d ) // same, using flags==0 default
```

and so we now have a nice high-level way of authoring APIs that can be consumed by anything that understands C, plus automates the annotation distinguishing between `inout` and `out` parameters and enforces it to be now always correct (never gets out of sync with the source because it's language-enforced).

Bibliography

[[Brooks 1975](#)] F. Brooks. *The Mythical Man-Month* (Addison-Wesley, 1975).

[[ECMA-334](#)] *Standard ECMA-334, C# Language Specification, 5th edition* (ECMA, 2017).

[[lib-ext 2017-02-08](#)] Discussion thread regarding making `std::move` implicit on last use of a path (WG21 email lists, 2017).

[[P1045R1](#)] D. Stone. “constexpr Function Parameters” (WG21 paper, 2019-09-27).

[[P1179R1](#)] H. Sutter. “Lifetime safety: Preventing common dangling (version 1.1)” (WG21 paper, 2019-11-22).

[Stroustrup 1994] B. Stroustrup. *The Design and Evolution of C++* (Addison-Wesley, 1994).

[[Sutter 2014](#)] H. Sutter. “Back to the basics! Essentials of modern C++ style” (CppCon, 2014).